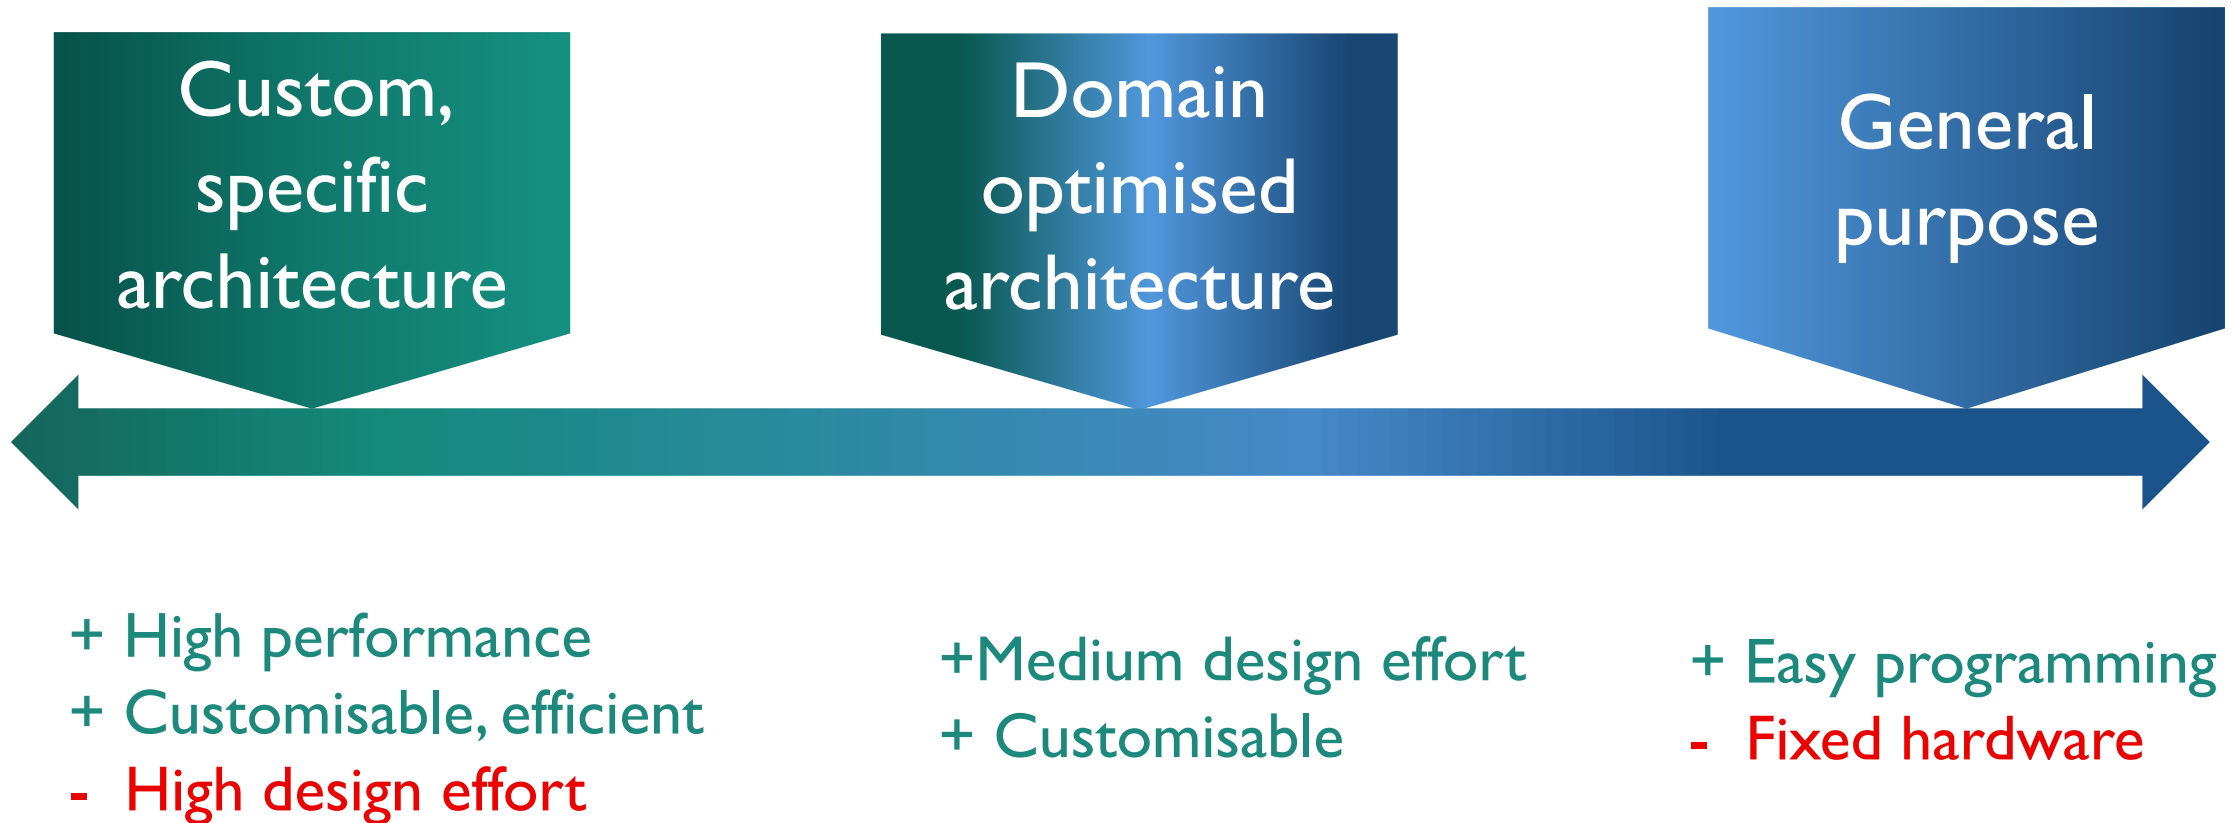


# Accelerating BFS Algorithm on a RISC-V-based Many-Core Cluster

**M. R. Ashuthosh<sup>1</sup>**, Aakarsh Vinay<sup>1</sup>, Krishna K. Nagar<sup>2</sup>, Madhura Purnaprajna<sup>1</sup>

<sup>1</sup>Centre for Heterogenous and Intelligent Processing System, PES University, Bengaluru, India

<sup>2</sup>Intel Corporation, USA



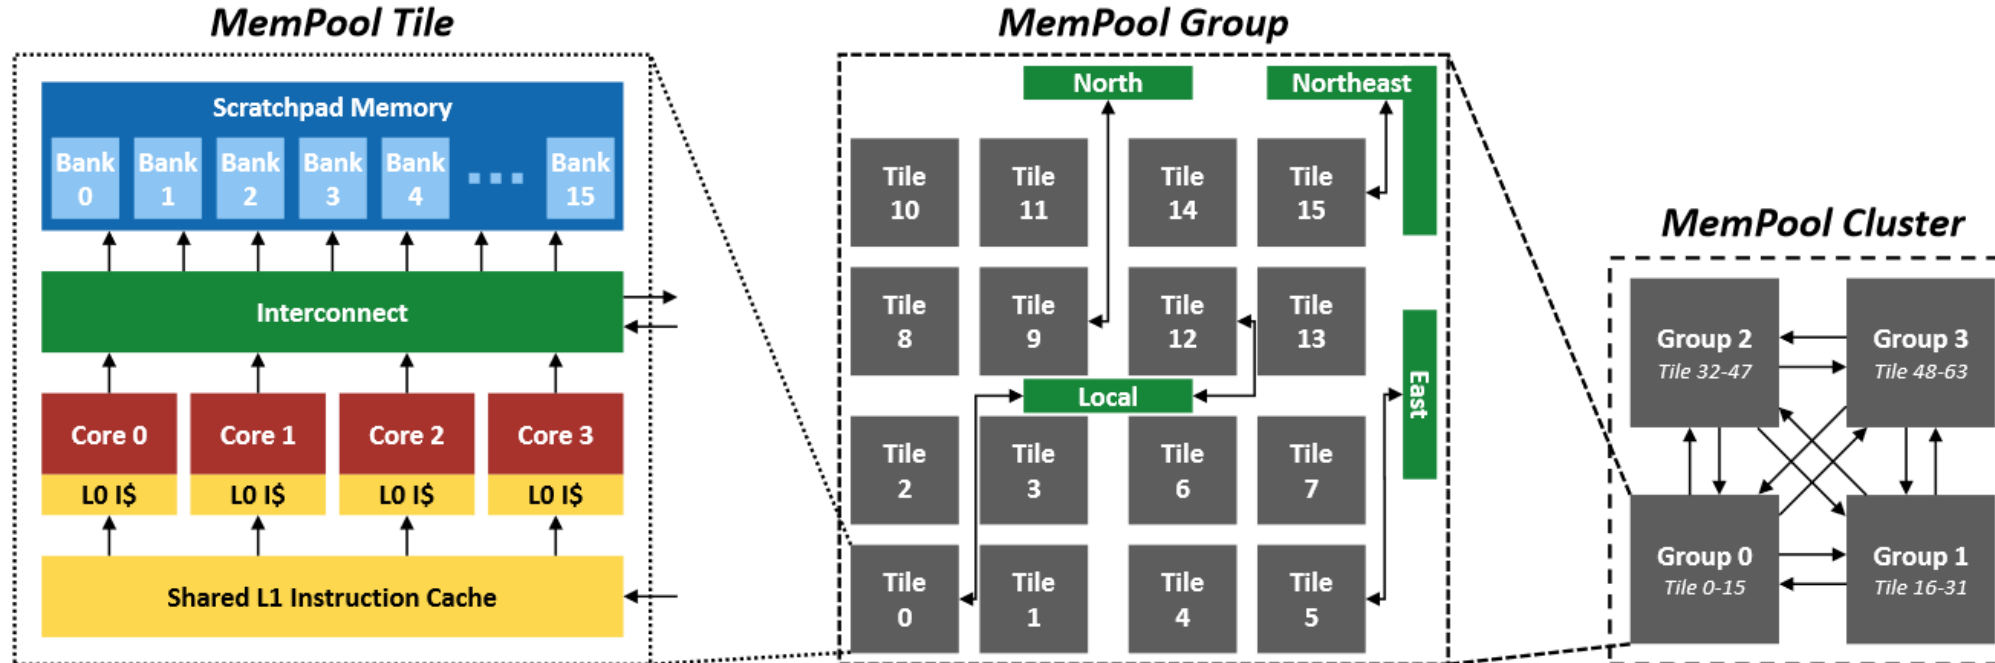
# Acceleration of graph algorithms

- Graph algorithms are everywhere



- Memory bound, irregular memory accesses
- BFS traversal is the basic component in graph algorithms

# Mempool [1]: A RISC-V-based Manycore Cluster



## Mempool Tile:

- 4 cores
- 16 memory banks
- Single-cycle latency

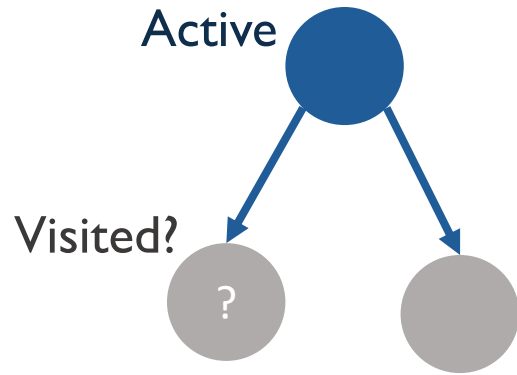
## Mempool Group:

- 64 cores
- 256 memory banks
- 3 cycles latency

## Mempool Cluster:

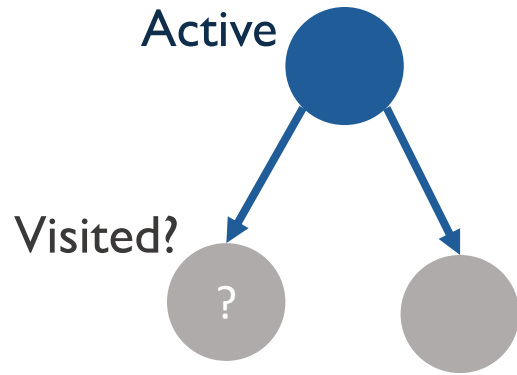
- 256 cores
- 1024 memory banks (1MiB)
- 5 cycles latency

## I. Top-down BFS



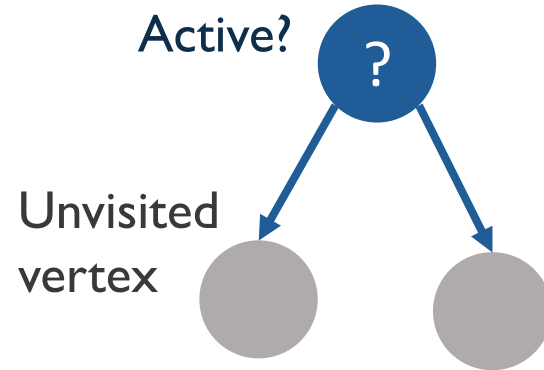
- Check neighbors of active vertex, if unvisited, update the unvisited vertex.

## 1. Top-down BFS



- Check neighbors of active vertex, if unvisited, update the unvisited vertex.

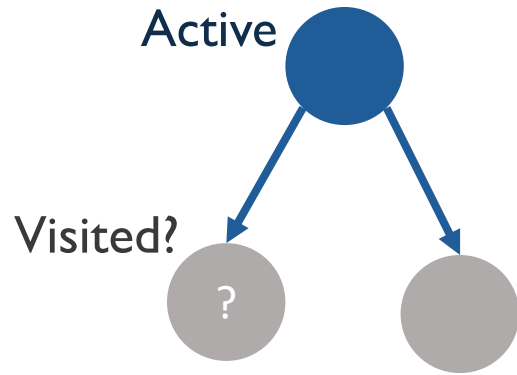
## 2. Bottom-up BFS



- Check neighbors of unvisited vertex, if active, update the unvisited vertex.

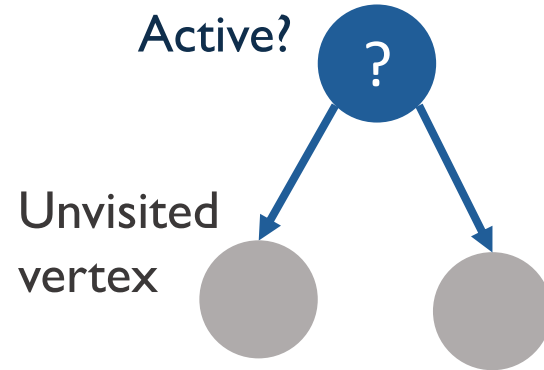
# BFS approaches

## 1. Top-down BFS



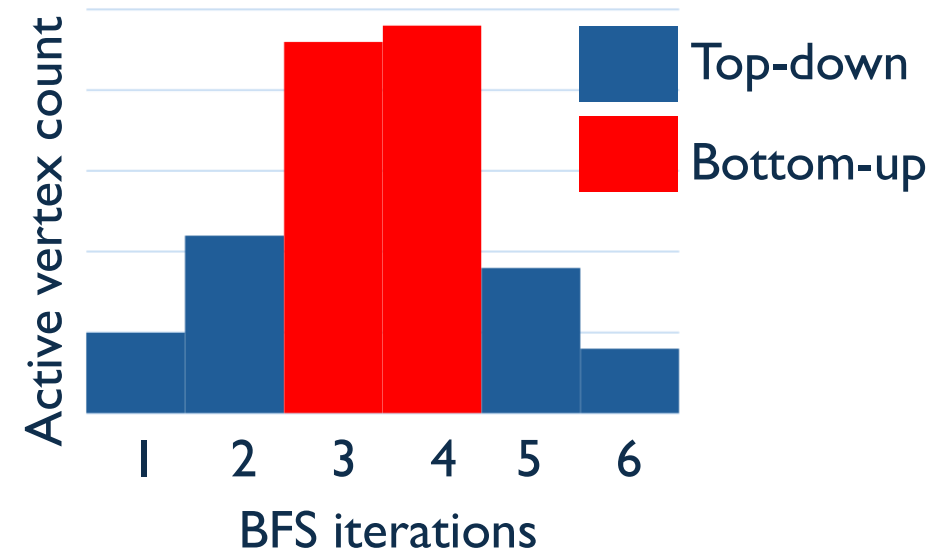
- Check neighbors of active vertex, if unvisited, update the unvisited vertex.

## 2. Bottom-up BFS

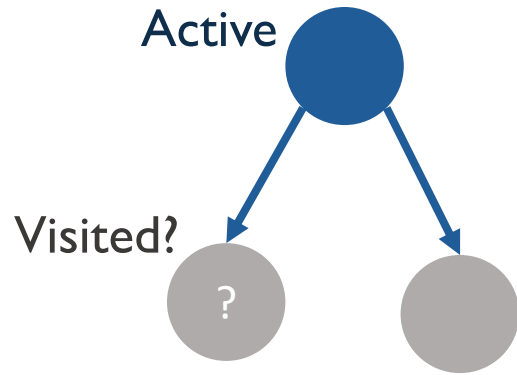


- Check neighbors of unvisited vertex, if active, update the unvisited vertex.

## 3. Hybrid BFS [3], an optimal approach

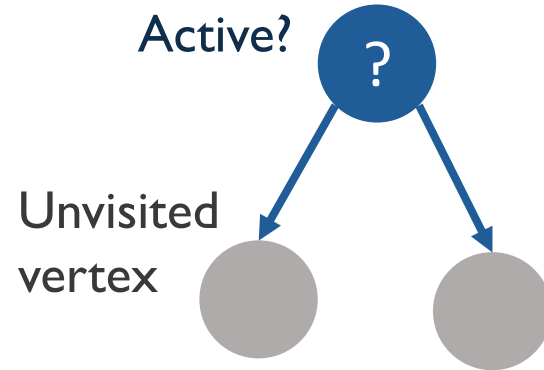


## 1. Top-down BFS



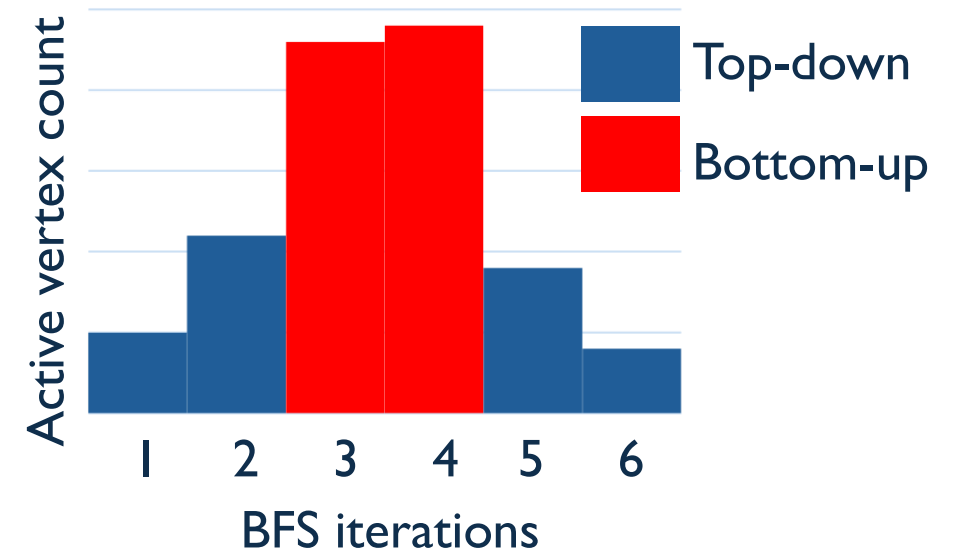
- Check neighbors of active vertex, if unvisited, update the unvisited vertex.

## 2. Bottom-up BFS



- Check neighbors of unvisited vertex, if active, update the unvisited vertex.

## 3. Hybrid BFS [3], an optimal approach

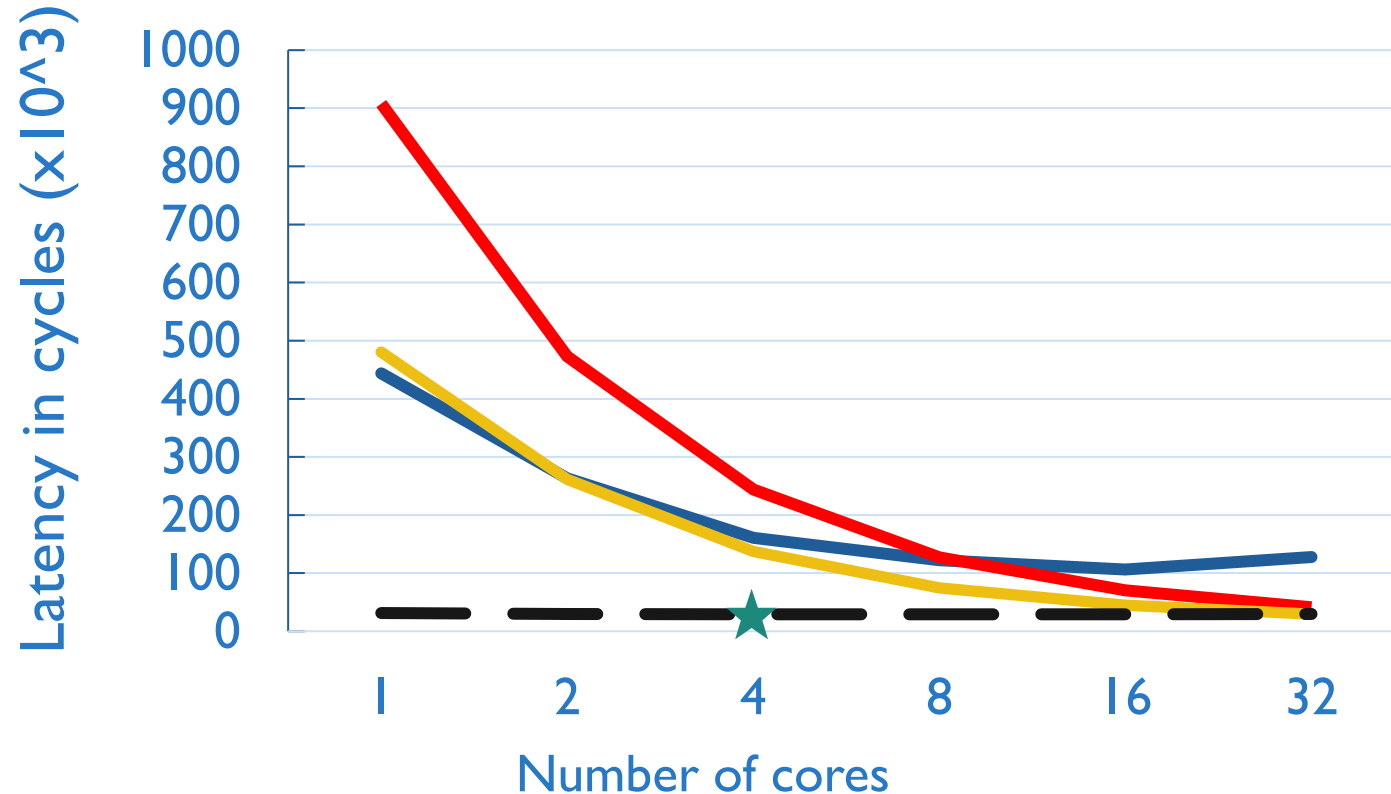


## 4. Custom hybrid BFS (our approach): Hybrid with optimal number of cores for iterations of top-down and bottom-up.



# How many cores are enough?

— Top-down BFS 
 — Bottom-up BFS 
 — Hybrid BFS 
 — Custom hybrid BFS

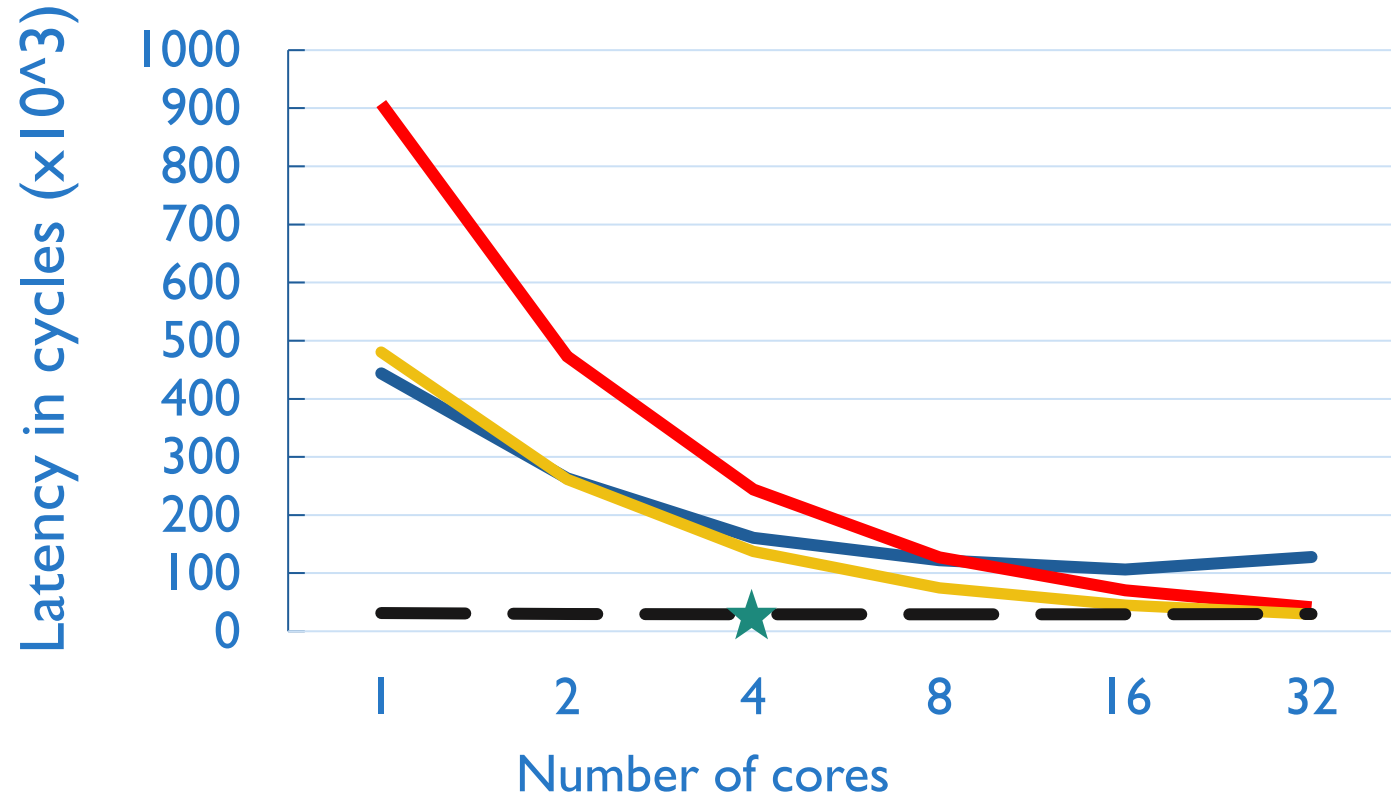


On traversing the tech-routers-rf [2], the speedups:

- Top-down BFS: 4.2× on 16 cores.
- Bottom-up BFS: 22.1× on 32 cores.
- Hybrid BFS: 16.02× on 32 cores.
- ★ Custom hybrid BFS: 16.5× when top-down=4 cores, bottom-up=32 cores.

# How many cores are enough?

— Top-down BFS — Bottom-up BFS — Hybrid BFS — Custom hybrid BFS



On traversing the tech-routers-rf [2], the speedups:

- Top-down BFS: 4.2× on 16 cores.
- Bottom-up BFS: 22.1× on 32 cores.
- Hybrid BFS: 16.02× on 32 cores.
- ★ Custom hybrid BFS: 16.5× when top-down=4 cores, bottom-up=32 cores.
- Varying the number of cores for top-down and fixing 32 cores for bottom-up provides no more than 7.4% performance improvement.

# Conclusion & Ongoing efforts

---

- Custom hybrid BFS(our approach) performs better than Hybrid BFS with maximum number of cores.
- These results underscore the importance of performance modeling for many-core architectures in rapid design-space exploration.
- Similar analysis for popular graph algorithms such as Pagerank.



Semiconductor  
Research  
Corporation



Scan to reach us!



**Thank you**

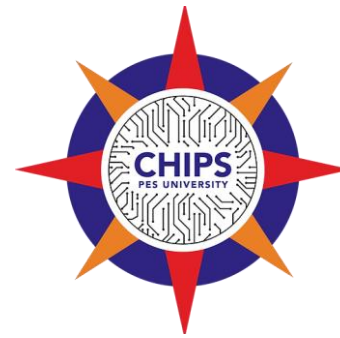
**Contact:**

Mail: [ashuthoshmr@pes.edu](mailto:ashuthoshmr@pes.edu)

Website: <https://ashuthosh.de/>

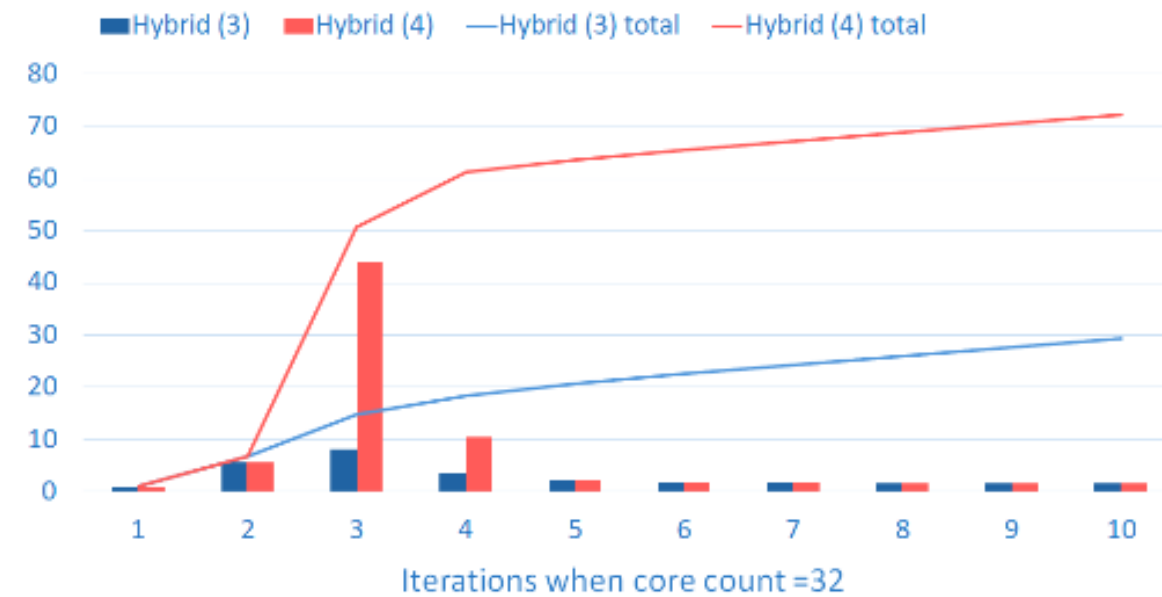
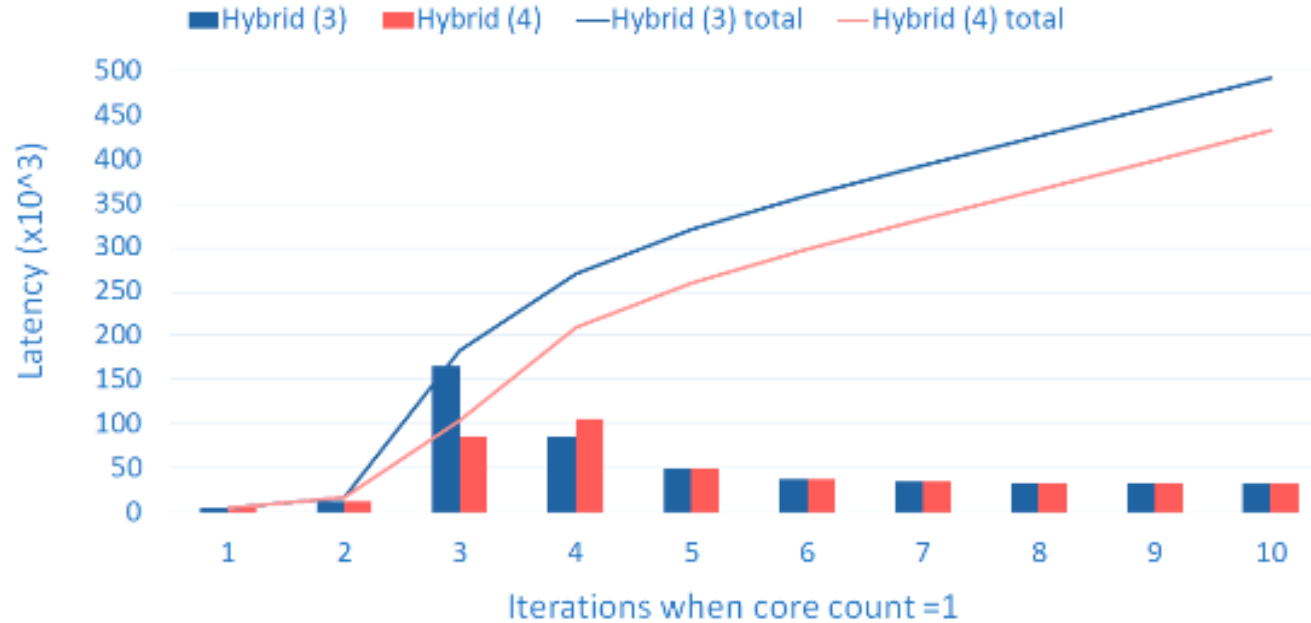


Semiconductor  
Research  
Corporation



## Backup

# Why hybrid is slower than top-down when #core=1



Why?	Edges explored	Edges explored
Core count	Switch at 3 (bu at 3)	Switch at 4 (td at 3)
1	6716	1964
2	3423	1010
4	1731	612
8	916	298
16	540	438
32	293	721